

Deepak Katukuri

Potsdam, Germany
deepakkatukuri@gmail.com | +49 1629741795

24 September 2026

SAP SE
BTP Core Team, Kyma Runtime
Garching bei München, Germany

Re: Working Student, AI Features for SAP BTP Kyma Runtime

Dear Hiring Team,

I'm applying for the Working Student role building AI features for Kyma runtime. I'm an MSc Data Science student based near Berlin with 2.5 years of production software engineering behind me, and most of my recent project work sits close to what this role does day to day: multi-agent systems, LangGraph, and Kubernetes.

I built JobPilot, a 16-node LangGraph pipeline that automates job-search CV tailoring end to end. It pulls evidence from a Chroma-backed store, refuses any CV claim it can't support with a source, checkpoints so a crash doesn't lose a run, and is held honest by 238 pytest tests across the codebase. It's the same category of problem this role names: agent skills, tools, and the automation layer around them, just applied to a different domain. I've also built and run my own MCP server, so the protocol side of "new skills and tools for agents" is familiar ground. A2A specifically I haven't worked with yet, and I'd rather say that directly than talk around it.

On Kubernetes: for a three-person cloud-computing project called CloudNotes, I owned the full K8s and AWS layer myself. That meant the Deployment with resource limits, the Service, Traefik Ingress, readiness and liveness probes, autoscaling from 2 to 5 pods at 60% CPU, and Secrets. I compared EKS against K3s on cost before picking K3s, and I built the GitHub Actions pipeline that ships it: multi-arch build, push to Docker Hub, SSH to the box, gated on kubectl rollout status. Since Kyma is a Kubernetes runtime, this is close to the actual tradeoffs your team is making, just at smaller scale.

The rest of my background is RAG. I hand-built a hybrid retriever, then validated it against a LangChain rebuild and caught a tokenizer regression that quietly cost 0.06 recall. I later added a corrective RAG loop on top, where the system re-retrieves on a weak match and refuses to answer when the corpus genuinely doesn't cover the question. Prompt design, context injection, and retrieval are all things I've built from the ground up, not just called through an API.

I'm currently in Potsdam and open to relocating to Munich for the weekly office day in Garching. I'm available now as a Werkstudent and full-time once I finish my MSc.

Thank you for considering my application. I'd welcome a conversation about any of this.

Best regards,
Deepak Katukuri